

OpenCV

February 14, 2019

1 OpenCV

Prerequisite: Before starting this exercise, you should make yourself familiar with Python and some necessary library, e.g., numpy, matplotlib, etc. One good tutorial can be found [here](#).

In this exercise you will:

- * Learn about some basic image processing operations with OpenCV.
- * Re-implement some basic image processing operations. This will help you to
- * Have better understand about the image processing operations.
- * Practice Python programming with Numpy library.

```
In [1]: import cv2
import numpy as np
import sys
import matplotlib
from matplotlib import pyplot as plt

# This is a bit of magic to make matplotlib figures appear inline in the notebook
# rather than in a new window.
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

In [2]: img = cv2.imread('imgs/grass.jpg')
img = cv2.resize(img, (100, 100))
print(img[:, :, 0])
img[:, :, np.array([0, 1, 2])] = img[:, :, np.array([2, 1, 0])]
print(img[:, :, 2])

[[16  7 17 ... 44 79 38]
 [28 28 30 ... 23 35 37]
 [22 24 34 ... 49 34 26]
 ...
 [86 49 24 ... 35 44 10]
 [30 30 31 ... 48 30 34]
 [23 25 41 ... 50 35 34]]
[[16  7 17 ... 44 79 38]
 [28 28 30 ... 23 35 37]
```

```
[22 24 34 ... 49 34 26]
...
[86 49 24 ... 35 44 10]
[30 30 31 ... 48 30 34]
[23 25 41 ... 50 35 34]]
```

```
In [3]: def rel_error(out, correct_out):
        return np.sum(abs(out.astype(np.float32) - correct_out.astype(np.float32)) /
                        (abs(out.astype(np.float32)) + abs(correct_out.astype(np.float32))))
```

```
In [4]: # Checking OpenCV version
        cv2.__version__
```

```
Out[4]: '4.0.0'
```

1.0.1 NOTICE:

In this lab exercise, we recommend to use OpenCV 3.x version, the documentations for OpenCV API can be found [here](#).

1.1 Load images

Use the function `cv2.imread()` to read an image. The image should be in the working directory or a full path of image should be given. The function will return a numpy matrix.

Second argument is a flag which specifies the way image should be read.

- `cv2.IMREAD_COLOR` - (1): Loads a color image. Any transparency (alpha channel) of image will be neglected. It is the **default flag**.
- `cv2.IMREAD_GRAYSCALE` - (0): Loads image in grayscale mode
- `cv2.IMREAD_UNCHANGED` - (-1): Loads image as such including alpha channel, if included.

NOTE: Color image loaded by OpenCV is in *Blue-Green-Red (BGR)* mode. But Matplotlib displays in *RGB* mode. So color images will not be displayed correctly in Matplotlib if image is read with OpenCV. We will discuss how to handle to display properly later.


```
In [5]: img_gray = cv2.imread('imgs/opencv_logo.png', 0)

plt.figure(figsize=(20,10))
plt.subplot(131),
plt.imshow(img_gray, cmap='gray') # include cmap='gray' to display gray image
plt.title('Gray'), plt.xticks([]), plt.yticks([])

img_color1= cv2.imread('imgs/opencv_logo.png', 1)
plt.subplot(132), plt.imshow(img_color1),
plt.title('Color'), plt.xticks([]), plt.yticks([])

img_color2= cv2.imread('imgs/opencv_logo.png', -1)
```

```
plt.subplot(133),plt.imshow(img_color2),
plt.title('Color'),plt.xticks([]), plt.yticks([])
plt.show()
```



Question: How many channels for each image: `img_gray`, `img_color1`, `img_color2`?

Your answer: * `img_gray`: * `img_color1`: * `img_color2`:

2 Transformations

2.1 Scaling

Resize image using the function `cv2.resize`.

In [6]: # Get list of available flags

```
flags = [i for i in dir(cv2) if i.startswith('INTER_')]
print(flags)
```

```
['INTER_AREA', 'INTER_BITS', 'INTER_BITS2', 'INTER_CUBIC', 'INTER_LANCZOS4', 'INTER_LINEAR', 'INTER_NEAREST', 'INTER_SPLINE']
```

In [7]: `img = cv2.imread('imgs/opencv_logo1.png', 1)`

```
res = cv2.resize(img, None, fx=2.0, fy=2.0, interpolation = cv2.INTER_CUBIC)
```

#OR

```
height, width = img.shape[:2]
```

```
res = cv2.resize(img, (2*width, 2*height), interpolation = cv2.INTER_CUBIC)
```

```
#####
# TODO: Check the size of 'img' and 'res'?
```

```
#####
print('Size of img:{}'.format(img.shape))
print('Size of res:{}'.format(res.shape))
#####
```

```

#                                     END OF YOUR CODE                                     #
#####

#####
# TODO: Resize 'img' so as to the smaller side is 500, while keeping image
# ratio unchanged.
#####
aspect_ratio = img.shape[1]/img.shape[0]
newimg = cv2.resize(img, (int(500*aspect_ratio),500), interpolation=cv2.INTER_CUBIC)
print('Size of resizedimg:{}'.format(newimg.shape))
#####
#                                     END OF YOUR CODE                                     #
#####

```

```

Size of img:(378, 428, 3)
Size of res:(756, 856, 3)
Size of resizedimg:(500, 566, 3)

```

2.2 Translation

Translation is the shifting of object's location. If you know the shift in (x, y) direction, let it be (t_x, t_y) , you can create the transformation matrix M as follows:

$$M = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix}$$

You can take make it into a Numpy array of type **np.float32** and pass it into **cv2.warpAffine()** function.

```

In [8]: img = cv2.imread('imgs/opencv_logo1.png', 1)
        rows,cols,_ = img.shape
        M = np.float32([[1,0,100],[0,1,50]]) # Shift right by 100 and down by 50
        dst = cv2.warpAffine(img,M,(cols,rows))
        print(res.shape)
#####
# TODO: Observed that the bottom right of 'dst' image is lost. Modifying the
# following codeline so as to the 'res' image is fully shown.
#####
resizedres = cv2.warpAffine(res, M, dsize = (956,806))
#####
#                                     END OF YOUR CODE                                     #
#####

plt.figure(figsize=(20,10))
plt.subplot(131),plt.imshow(img),
plt.title('Original'),plt.xticks([], plt.yticks([]))
plt.subplot(132),plt.imshow(dst),
plt.title('Shifted images'),plt.xticks([], plt.yticks([]))

```

```
plt.subplot(133),plt.imshow(resizedres),
plt.title('Shifted images'),plt.xticks([]), plt.yticks([])
plt.show()
```

(756, 856, 3)



2.3 Rotation

Calculates an affine matrix of 2D rotation using `cv2.getRotationMatrix2D()`. * 1st argument: center
* 2nd argument: angle (in degree) * 3rd argument: scale

```
In [9]: img = cv2.imread('imgs/opencv_logo1.png', 1)
        H,W,_ = img.shape
        #####
        # TODO: Run the code to observe the output image.
        # Modify the code below so as to the 'dst' image has no black padding.
        #####
        adjusted = cv2.getRotationMatrix2D((H/1.76,W/2), 90, 1)
        newimg = cv2.warpAffine(img, adjusted, (H,W))
        #####
        #                                     END OF YOUR CODE                                     #
        #####

        plt.imshow(newimg),
        plt.title('Rotated images'),plt.xticks([]), plt.yticks([])
        plt.show()
```




3 Changing color space - Grayscale

Grayscale values is converted from RGB values by a weighted sum of the R, G, and B components:

$$0.299 \times R + 0.587 \times G + 0.114 \times B$$

```
In [10]: # Split channels
img = cv2.imread('imgs/balls.jpg', 1)

plt.figure(figsize=(20,10))
plt.subplot(131),plt.imshow(img[:, :,0], cmap='gray'),
plt.title('Blue channel'),plt.xticks([]), plt.yticks([])
plt.subplot(132),plt.imshow(img[:, :,1], cmap='gray'),
plt.title('Green channel'),plt.xticks([]), plt.yticks([])
plt.subplot(133),plt.imshow(img[:, :,2], cmap='gray'),
plt.title('Red channel'),plt.xticks([]), plt.yticks([])
plt.show()
```



```

In [11]: def rgb2gray(img):
        """
        A implementation of the method that converts BGR image to grayscale image of
        uint8 data type.
        """
        out = 0.114*img[:, :, 0] + 0.587*img[:, :, 1] + 0.299*img[:, :, 2]
        #####
        # TODO: Implement the method to convert BGR image to Grayscale image.          #
        # Hint: Remember to round and convert the values to nearest uint8 values.      #
        #####
        out = np.round(out)
        #####
        #                                     END OF YOUR CODE                          #
        #####

        return out

```

Run the following code section to compare your implementation of the `rgb2gray` function with OpenCV built-in function `cv2.cvtColor`.


```

In [12]: img = cv2.imread('imgs/ball_red.jpg', 1)
        img_gray1 = rgb2gray(img)
        img_gray2 = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

        plt.figure(figsize=(20,10))
        plt.subplot(121),plt.imshow(img_gray1, cmap='gray'),
        plt.title('My rgb2gray'),plt.xticks([]), plt.yticks([])
        plt.subplot(122),plt.imshow(img_gray1 - img_gray2, cmap='gray'),
        plt.title('Difference'),plt.xticks([]), plt.yticks([])
        plt.show()

        # Check your output: count
        print('Testing rgb2gray')
        print('Number of difference pixel is %d' % np.count_nonzero(img_gray1 - img_gray2))

```



Testing rgb2gray
Number of difference pixel is 1

Question: Does your implementation of rgb2gray function give the result that is exactly the same as OpenCV built-in function? Why?

Your answer: fill in here

No it does not. It differs by 1 due to numeric calculation differences

4 Changing color space - Detect object by color.

By converting BGR image to HSV, we can use this to extract a colored object. In HSV, it is more easier to represent a color than RGB color-space. In this exercise, we will try to extract blue, red, and yellow colored objects. So here is the method:

- Take each frame of the video
- Convert from BGR to HSV color-space
- We threshold the HSV image for a range of blue color
- Now extract the blue object alone, we can do whatever on that image we want.

```
In [13]: # Get list of available flags
         flags = [i for i in dir(cv2) if i.startswith('COLOR_')]
         print (flags)
```

```
['COLOR_BAYER_BG2BGR', 'COLOR_BAYER_BG2BGR_A', 'COLOR_BAYER_BG2BGR_EA', 'COLOR_BAYER_BG2BGR_VNG
```

```
In [14]: frame = cv2.imread('imgs/balls.jpg', 1)
```



```

# Convert BGR to RGB, now you will see the color of 'frame' image
# is displayed properly.
frame = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)

# Convert BGR to HSV
hsv = cv2.cvtColor(frame, cv2.COLOR_RGB2HSV)

# define range of blue color in HSV
lower_blue = np.array([110,50,50])
upper_blue = np.array([130,255,255])

# Threshold the HSV image to get only blue colors
mask = cv2.inRange(hsv, lower_blue, upper_blue)

# Bitwise-AND mask and original image
res = cv2.bitwise_and(frame, frame, mask= mask)


#####
# TODO: Implement masks for red and yellow balls.
#####
lower_red = np.array([0, 100, 100])
upper_red = np.array([255, 255, 255])
red_mask = cv2.inRange(hsv, lower_red, upper_red)
res = cv2.bitwise_and(frame, frame, mask= red_mask)

# lower_green = np.array([50,110,50])
# upper_green = np.array([255,130,255])
# green_mask = cv2.inRange(hsv, lower_green, upper_green)
# res = cv2.bitwise_and(frame, frame, mask= green_mask)

#####
#                                     END OF YOUR CODE                                     #
#####

plt.figure(figsize=(20,10))
plt.subplot(131),plt.imshow(frame),
plt.title('Original'),plt.xticks([]), plt.yticks([])
plt.subplot(132),plt.imshow(red_mask, cmap='gray'),
plt.title('Mask'),plt.xticks([]), plt.yticks([])
plt.subplot(133),plt.imshow(res),
plt.title('Output'),plt.xticks([]), plt.yticks([])
plt.show()

```



5 2D Convolution (Image Filtering)

OpenCV provides a function, `cv2.filter2D`, to convolve a kernel with an image.

```
In [15]: def convolution_naive(x, F, conv_param):
         """
         A naive implementation of a convolutional filter.

         The input consists of a gray scale image x (1 channel) with height H and width
         W. We convolve each input with filter F, which has height HH and width HH.

         Input:
         - x: Input data of shape (H, W)
         - F: Filter weights of shape (HH, WW)
         - conv_param: A dictionary with the following keys:
           - 'stride': The number of pixels between adjacent receptive fields in the
             horizontal and vertical directions.
           - 'pad': The number of pixels that will be used to zero-pad the input.

         Return:
         - out: Output data, of shape (H', W') where H' and W' are given by
            $H' = 1 + (H + 2 * pad - HH) / stride$ 
            $W' = 1 + (W + 2 * pad - WW) / stride$ 
         """

         stride = conv_param['stride']
         pad = conv_param['pad']

         H, W = x.shape
         HH, WW = F.shape
         frame = (WW - 1) // 2
         H_prime = int(1 + (H + 2 * pad - HH) / stride)
         W_prime = int(1 + (W + 2 * pad - WW) / stride)
         #print(H_prime, W_prime)
         x_pad = np.lib.pad(x, ((pad, pad), (pad, pad)),\
                             'constant', constant_values=(0))
         out = np.zeros((H_prime, W_prime), dtype=x.dtype)
         print(x_pad.shape)
```

```
#####
# TODO: Implement the convolutional forward pass. #
# Hint: Using 2 nested for-loop to calculate each pixel of the output image. #
#####
for i in range(pad, H+pad):
    for j in range(pad, W+pad):
        roi = x_pad[i - pad:i + pad + 1, j - pad:j + pad + 1]
        k = (roi * F).sum()
        out[i - pad, j - pad] = k
#####
#                                     END OF YOUR CODE                                     #
#####
return out
```

Run the following code section to test your implementation of the convolution_naive function

```
In [16]: x_shape = (5, 5)
         F_shape = (3, 3)
         x = np.linspace(-0.1, 0.5, num=np.prod(x_shape)).reshape(x_shape)
         F = np.linspace(-0.2, 0.3, num=np.prod(F_shape)).reshape(F_shape)
         conv_param = {'stride': 1, 'pad': 1}

         out = convolution_naive(x, F, conv_param)
         correct_out = np.array( [[ 0.0075,      0.030625,   0.0521875,  0.07375,   0.0475   ]
                                [ 0.114375,   0.1725,     0.18375,   0.195,     0.10875  ]
                                [ 0.1753125,  0.22875,     0.24,       0.25125,   0.1228125]
                                [ 0.23625,    0.285,       0.29625,   0.3075,    0.136875 ]
                                [ 0.0075,    -0.05375,   -0.0603125, -0.066875, -0.1025   ]
                                )
         # print(correct_out.shape)
         # print(out)

         # Compare your output to ours; difference should be very small
         print('Testing convolution_naive')
         print('difference: ', rel_error(out, correct_out))

(7, 7)
Testing convolution_naive
difference:  0.0

In [17]: # List of available BORDER effect
         flags = [i for i in dir(cv2) if i.startswith('BORDER_')]
         print (flags)

['BORDER_CONSTANT', 'BORDER_DEFAULT', 'BORDER_ISOLATED', 'BORDER_REFLECT', 'BORDER_REFLECT101']
```

This is done by convolving image with a normalized box filter. A 5×5 normalized box filter would look like below:

```
In [18]: # Convert image data type from uint8 to float32.
img = cv2.imread('imgs/text.png', 1).astype(np.float32)
img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

kernel = np.zeros((5,5), np.float32)
#####
# TODO: Create a 5x5 kernel as K shown above. #
#####
kernel = np.ones((5,5), np.float32)/25
#####
#                                     END OF YOUR CODE #
#####
blur_2dfilter = cv2.filter2D(img,-1,kernel)

# The above codes can be replaced by the following code line.
blur = cv2.blur(img,(5,5))

# Check your output; difference should be around 4e-3
print('Testing convolution_naive')
print('difference: ', rel_error(blur_2dfilter, blur))

# Visualize the output image
plt.figure(figsize=(20,10))
plt.subplot(121),plt.imshow(img, cmap='gray'),
plt.title('Original'),plt.xticks([]), plt.yticks([])
plt.subplot(122),plt.imshow(blur, cmap='gray'),
plt.title('Average Blur'),plt.xticks([]), plt.yticks([])
plt.show()
```

12